

Tools Seminar

Week 8 - Deep Learning

Hongzheng Chen

Mar 16, 2020

- 1 Introduction
- 2 Deep Learning
- 3 Frameworks
 - Installation
 - Tutorials
- 4 Summary

1

Introduction

AI Milestones

- 1950, Alan Turing proposed the famous Turing test

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)
- 1997, IBM’s Deep Blue beat world chess champion Garry Kasparov

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)
- 1997, IBM’s Deep Blue beat world chess champion Garry Kasparov
- 2011, IBM Watson defeated two champions at quiz show Jeopardy

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)
- 1997, IBM’s Deep Blue beat world chess champion Garry Kasparov
- 2011, IBM Watson defeated two champions at quiz show Jeopardy
- 2012, Jeff Dean and Andrew Ng used unsupervised learning to train neural network which learned to recognize cats

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)
- 2016, Google DeepMind's AlphaGo defeated 9-dan Go master Lee sedol by 4:1

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)
- 2016, Google DeepMind's AlphaGo defeated 9-dan Go master Lee sedol by 4:1
 - Firstly proposed deep reinforcement learning

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in **ImageNet Large Scale Visual Recognition Challenge**
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)
- 2016, Google DeepMind's AlphaGo defeated 9-dan Go master Lee sedol by 4:1
 - Firstly proposed deep reinforcement learning
- 2018, Google **BERT** model achieved the state-of-the-art performance in 11 NLP tasks

ImageNet & Deep Neural Network

ImageNet [Feifei Li, Stanford] Large Scale Visual Recognition Challenge

	LeNet	AlexNet	VGG	GoogleLeNet	ResNet
# of Layers	5	8	19	22	152
Top 5 Error	N/A	16.4%	7.3%	6.7%	3.57%
Year	1994	2012	2014	2014	2015

This is why it's called “**deep**” learning

2018 Turing Award

2018 Turing Award: Geoffrey Hinton, Yoshua Bengio, Yann LeCun

“for conceptual and engineering breakthroughs that have made deep neural networks a critical component of computing”

- Geoffrey Hinton: Backpropagation, Boltzmann Machines, Improvements to Convolutional Neural Network (CNN)
- Yoshua Bengio: Probabilistic models of sequences, High-dimensional word embeddings and attention, Generative adversarial networks (GAN)
- Yann LeCun: CNN, backprop, Broadening the vision of neural networks


Impetus of Deep Learning

Looking back, we may know what leads to the boom of DL in 2010s

- Large amount of labeled **data**: ImageNet
- Improvement of **hardware**: GPU \rightarrow GPGPU (general-purpose GPU)
- Improvement of **algorithms**: deep networks, dropout

Introductory Books and Courses

- Feifei Li, [Stanford cs231n](#): Convolutional Neural Networks for Visual Recognition (highly recommended!)
- Chris Manning, [Stanford cs224n](#): Natural Language Processing with Deep Learning
- Ian Goodfellow, [Deep Learning Chinese version](#)



2

Deep Learning

Linear Regression


Recall the linear regression problem

$$y = \mathbf{w}^T \mathbf{x} + b = \begin{bmatrix} \mathbf{x}^T & b \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix} = \boldsymbol{\theta}^T \mathbf{x}$$


To minimize loss function (MSE)

$$\min_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = \frac{1}{m} \sum_{i=1}^m (y_i - \mathbf{w}^T \mathbf{x}_i - b)^2$$


But linear function can only deal with linearly separable problems



x_1 and x_2



x_1 or x_2



x_1 xor x_2

Powerful Models

Can we build a model that is powerful enough to represent all the functions?

- $f(\text{image}) = \text{location}$
- $f(\text{question}) = \text{answer}$
- $f(\text{speech}) = \text{text}$
- ...

Powerful Models


Can we build a model that is powerful enough to represent all the functions?

- $f(\text{image}) = \text{location}$
- $f(\text{question}) = \text{answer}$
- $f(\text{speech}) = \text{text}$
- ...

Actually, in ML area, we have built decision tree, SVM, etc., but they usually need feature engineering and are not flexible

Neuron Model

We can refer to our brain and see how we learn




A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Fig source: <http://cs231n.github.io/neural-networks-1/>

Neuron Model

We can refer to our brain and see how we learn




A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Fig source: <http://cs231n.github.io/neural-networks-1/>

A neuron is just a linear model with activation function!

Activation Function

Used to add non-linear part to the model, enabling it to approximate much more complex functions



- Sigmoid: $g(z) = 1/(1 + e^{-z})$ (S curve)
- Tanh: $g(z) = \tanh(z)$
- ReLU (Rectified Linear Unit): $g(z) = \max(0, z)$, can avoid gradient vanishing

From one to more

Only one neuron can do limited things, what about more?

- More neurons in width:

A feed-forward network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of \mathbb{R}^n .

— *The Universal Approximation Theorem*

The question is that the theorem does not tell us how many neurons we need

From one to more

Only one neuron can do limited things, what about more?

- More neurons in width:


A feed-forward network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of \mathbb{R}^n .

— *The Universal Approximation Theorem*

The question is that the theorem does not tell us how many neurons we need

- More neurons in depth: We have multi-layer perceptron (MLP) — the basic model of nowadays deep learning!

Multi-Layer Perceptron (MLP) / Fully-connected NN




- Input layer: \mathbf{x}
- Hidden layer: $h(\mathbf{x}) = g(\mathbf{z}^h(\mathbf{x})) = g(\mathbf{W}^h \mathbf{x} + \mathbf{b}^h)$
- Output layer: $f(\mathbf{x}) = \sigma(\mathbf{z}^o(\mathbf{x})) = \sigma(\mathbf{W}^o \mathbf{h}(\mathbf{x}) + \mathbf{b}^o)$

* Softmax function: change output to probability (K -dimensional)

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

Classifier Training

- Find optimal parameters θ^* of a classifier $y = f(x; \theta)$
- Rule: given input x , classifier output $f(x; \theta)$ should be as close to the ideal output as possible



Classifier Training

Use MSE or other loss function

$$\min_{\boldsymbol{\theta}} \frac{1}{N} \sum_{i=1}^N \|f(\mathbf{x}_i; \boldsymbol{\theta}) - y_i\|_2^2$$

Use gradient descent to optimize parameters

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} - \alpha \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta})$$

About how to optimize the above function on NN (backpropagation), please read <http://cs231n.github.io/optimization-1/>

Different Kinds of NN

- **CNN** (convolutional NN): CV
 - Pooling / subsampling
 - Dropout
 - Residual block
- RNN (recurrent NN): NLP
 - LSTM
 - GRU
- GAN (generative adversarial network): Image generation

3

Frameworks

Deep Learning Frameworks

Framework: A large package consisting of lots of deep learning primitives/operators, and users can easily call them by API

- Google: [Tensorflow](#) (commonly used in industry)
 - Static computation graph
 - Jeff Dean
- Facebook: [PyTorch](#) (commonly used in academics)
 - Dynamic computation graph
 - Yangqing Jia, Caffe
- Amazon: MXNet
 - Tianqi Chen (UW → CMU)

* We focus on **PyTorch** in this seminar

PyTorch

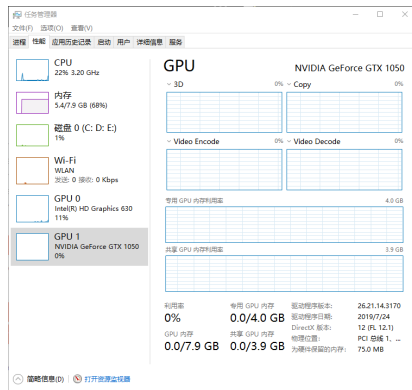
PyTorch: A Python-based scientific computing package

- A replacement for NumPy to use the power of GPUs
- A deep learning research platform that provides maximum flexibility and speed

Since it is highly embedded in Python, PyTorch is very Pythonic and easy-to-use

Pytorch Installation

Firstly check if your computer has discrete graphics card (GPU)



Install Nvidia driver: <https://zhuanlan.zhihu.com/p/54350088>

- CUDA 10.1
- cuDNN 7: [Installation guide](#)

Pytorch Installation

Select your configuration on this [website](#) and run the installation command

- Windows: Need to install Anaconda first
- WSL does not support GPU! Do NOT install Pytorch on WSL!
- Mac does not support GPU too (if you do not have external interface)!

e.g. For Windows with no GPUs

```
pip install torch==1.4.0+cpu torchvision==0.5.0+cpu -f https://download.pytorch.org/whl/torch_stable.html
```

Check if GPU works correctly by

```
import torch
print(torch.cuda.is_available())
```

Tutorials

PyTorch has very detailed documentations, make the best of them!

- Tutorials: <https://pytorch.org/tutorials/>
- Chinese tutorials: <https://pytorch.apachecn.org/>
- Documentation / API:
<https://pytorch.org/docs/stable/index.html>
- Deep Learning with PyTorch: A 60 Minute Blitz
 - Chinese version
 - You can download the .ipynb file or directly run on [Colab](#)

4

Summary

Summary

- Introduction
- Deep Learning Framework: PyTorch
 - Once you get in some trouble concerning PyTorch, you can search [the Docs of PyTorch](#) for details. Alternatively you can try to find if there are similar problems on [PyTorch Discuss](#).
- Get through cs231n!