

Tools Seminar

Week 8 - Deep Learning

Hongzheng Chen

Apr 5, 2020

- 1 Introduction
- 2 Deep Learning
- 3 Frameworks
 - Installation
 - Tutorials
- 4 Summary
- 5 Assignment

1

Introduction

AI Milestones

- 1950, Alan Turing proposed the famous Turing test

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)
- 1997, IBM’s Deep Blue beat world chess champion Garry Kasparov

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)
- 1997, IBM’s Deep Blue beat world chess champion Garry Kasparov
- 2011, IBM Watson defeated two champions at quiz show Jeopardy

AI Milestones

- 1950, Alan Turing proposed the famous Turing test
- 1955, John McCarthy created the term “artificial intelligence” (1971 Turing award)
- 1997, IBM’s Deep Blue beat world chess champion Garry Kasparov
- 2011, IBM Watson defeated two champions at quiz show Jeopardy
- 2012, Jeff Dean and Andrew Ng used unsupervised learning to train neural network which learned to recognize cats

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)
- 2016, Google DeepMind's AlphaGo defeated 9-dan Go master Lee sedol by 4:1

AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)
- 2016, Google DeepMind's AlphaGo defeated 9-dan Go master Lee sedol by 4:1
 - Firstly proposed deep reinforcement learning


AI Milestones (Cont.)

- 2012, AlexNet achieved an error rate of only 16% in [ImageNet Large Scale Visual Recognition Challenge](#)
 - Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton (University of Toronto)
 - 2 GPU, a week, halved the error rate
 - **Open the era of deep learning**
- 2015, Kaiming He (MSRA) proposed ResNet which made machine see better than human (3.6% error rate in ImageNet)
- 2016, Google DeepMind's AlphaGo defeated 9-dan Go master Lee sedol by 4:1
 - Firstly proposed deep reinforcement learning
- 2018, Google [BERT](#) model achieved the state-of-the-art performance in 11 NLP tasks

ImageNet & Deep Neural Network

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) [Feifei Li, Stanford]

- Labeled dataset → supervised learning
- 14+ million images, 20,000 categories



This is why it's called “**deep**” learning

2018 Turing Award

2018 Turing Award: Geoffrey Hinton, Yoshua Bengio, Yann LeCun

“for conceptual and engineering breakthroughs that have made deep neural networks a critical component of computing”

- Geoffrey Hinton: Backpropagation, Boltzmann Machines, Improvements to Convolutional Neural Network (CNN)
- Yoshua Bengio: Probabilistic models of sequences, High-dimensional word embeddings and attention, Generative adversarial networks (GAN)
- Yann LeCun: CNN, backprop, Broadening the vision of neural networks (LeNet5)

Impetus of Deep Learning


Looking back, we may know what leads to the boom of DL in 2010s

- Large amount of labeled **data**: ImageNet (2010)
- Improvement of **algorithms**: deep networks, dropout (2012)
- Invention of deep learning **systems**: TensorFlow (2015), PyTorch (2016)
- Improvement of **hardware**: GPU \rightarrow GPGPU (general-purpose GPU)

All of them are indispensable and make up the whole DL stack

Introductory Books and Courses

- Feifei Li, [Stanford cs231n](#): Convolutional Neural Networks for Visual Recognition (highly recommended!)
- Chris Manning, [Stanford cs224n](#): Natural Language Processing with Deep Learning
- Ian Goodfellow, [Deep Learning](#), Chinese version



2

Deep Learning

Linear Regression


Recall the linear regression problem

$$y = \mathbf{w}^T \mathbf{x} + b = \begin{bmatrix} \mathbf{x}^T & b \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix} = \boldsymbol{\theta}^T \mathbf{x}$$


To minimize loss function (MSE)

$$\min_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = \frac{1}{m} \sum_{i=1}^m (y_i - \mathbf{w}^T \mathbf{x}_i - b)^2$$


But linear function can only deal with linearly separable problems



x_1 and x_2



x_1 or x_2



x_1 xor x_2

Powerful Models

Can we build a model that is powerful enough to represent all the functions?

- $f(\text{image}) = \text{location}$
- $f(\text{question}) = \text{answer}$
- $f(\text{speech}) = \text{text}$
- ...

Powerful Models

Can we build a model that is powerful enough to represent all the functions?

- $f(\text{image}) = \text{location}$
- $f(\text{question}) = \text{answer}$
- $f(\text{speech}) = \text{text}$
- ...

Actually, in ML area, we have built decision tree, SVM, etc., but they usually need feature engineering and are not flexible

We need to reduce the burden of programmers
and make machines more automatic & intelligent

We need to reduce the burden of programmers
and make machines more automatic & intelligent

But how?

Representation Learning




Fig source: [Deep Learning](#) book

Representation matters!

Representation Learning

Different feature representation affects final performance

Representation Learning

Different feature representation affects final performance

Then, let machine learn the feature itself!

features \rightarrow mapping from features \rightarrow output

e.g.

- Encoder-decoder
- Word embeddings, graph embeddings

But...

We may have lots of features...

e.g. figure out what the object is in the photo

- size
- color
- material
- illumination
- view angle
- ...

Representation learning captures several features, but cannot capture all of them

But...

We may have lots of features...

e.g. figure out what the object is in the photo

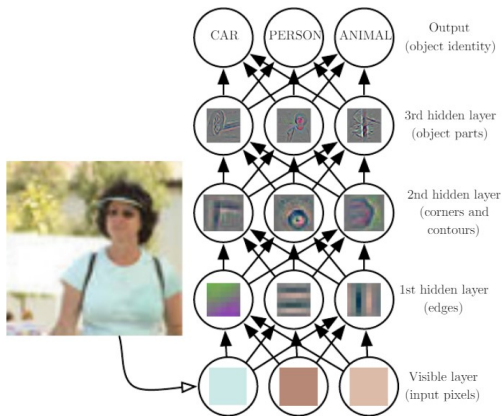
- size
- color
- material
- illumination
- view angle
- ...

Representation learning captures several features, but cannot capture all of them

So complex!

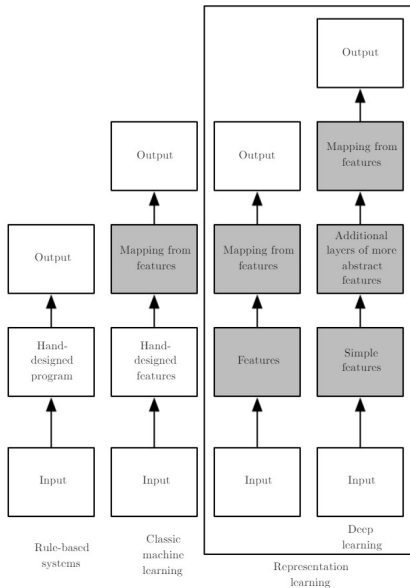
Decouple to simple features!

Learn from simple/shallow features and gradually to complex/deep features




Key: Get deeper!

Fig source: *Deep Learning* book

Fig source: *Deep Learning* book

Neuron Model

We can refer to our brain and see how we learn




A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Fig source: <http://cs231n.github.io/neural-networks-1/>

Neuron Model

We can refer to our brain and see how we learn




A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Fig source: <http://cs231n.github.io/neural-networks-1/>

A neuron is just a linear model with **activation function**!

Activation Function

Add non-linear part to the model, enabling it to approximate much more **complex** functions (key of NN!)



- Sigmoid: $g(z) = 1/(1 + e^{-z})$ (S curve)
- Tanh: $g(z) = \tanh(z)$
- ReLU (Rectified Linear Unit): $g(z) = \max(0, z)$, can avoid gradient vanishing

From one to more

Only one neuron can do limited things, what about more?

- More neurons in width:

A feed-forward network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of \mathbb{R}^n .

— *The Universal Approximation Theorem*

The question is that the theorem does not tell us how many neurons we need

From one to more

Only one neuron can do limited things, what about more?

- More neurons in width:


A feed-forward network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of \mathbb{R}^n .

— *The Universal Approximation Theorem*

The question is that the theorem does not tell us how many neurons we need

- More neurons in depth: We have multi-layer perceptron (MLP) — the basic model of nowadays deep learning!

Multi-Layer Perceptron (MLP) / Fully-connected NN




- Input layer: \mathbf{x}
- Hidden layer: $h(\mathbf{x}) = g(\mathbf{z}^h(\mathbf{x})) = g(W^h \mathbf{x} + \mathbf{b}^h)$
- Output layer: $f(\mathbf{x}) = \sigma(\mathbf{z}^o(\mathbf{x})) = \sigma(W^o h(\mathbf{x}) + \mathbf{b}^o)$

* Softmax function: change output to probability (K -dimensional)

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

Classifier Training

- Find optimal parameters θ^* of a classifier $y = f(x; \theta)$
- Rule: given input x , classifier output $f(x; \theta)$ should be as close to the ideal output as possible



Classifier Training

Use MSE or other loss function

$$\min_{\boldsymbol{\theta}} \frac{1}{N} \sum_{i=1}^N \|f(\mathbf{x}_i; \boldsymbol{\theta}) - y_i\|_2^2$$

Use gradient descent to optimize parameters

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} - \alpha \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta})$$

About how to optimize the above function on NN (backpropagation), please read <http://cs231n.github.io/optimization-1/>

Training

Let's see how NN trains:

Tensorflow Playground

Different Kinds of NN

- **CNN** (convolutional NN): CV
 - Pooling / subsampling
 - Dropout
 - Residual block
- RNN (recurrent NN): NLP
 - LSTM
 - GRU
- GAN (generative adversarial network): Image generation

3

Frameworks

Deep Learning Frameworks

Framework: A large package consisting of lots of deep learning primitives/operators, and users can easily call them by API

- Google: [Tensorflow](#) (commonly used in industry)
 - Static computation graph
 - Jeff Dean
 - Facebook: [PyTorch](#) (commonly used in academics)
 - Dynamic computation graph
 - Yangqing Jia, Caffe
 - Amazon: [MXNet](#)
 - Tianqi Chen
- [Domestic] [PaddlePaddle](#) (Baidu), [Mindspore](#) (Huawei), [MegEngine](#) (Face++), [Jittor](#) (Tsinghua)

* We focus on **PyTorch** in this seminar

PyTorch

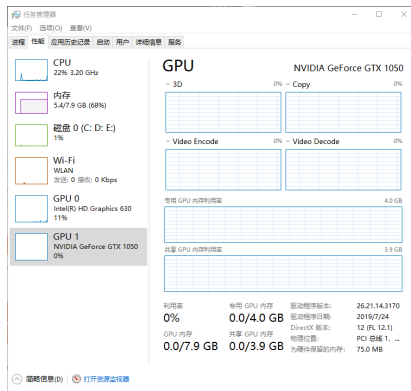
PyTorch: A Python-based deep learning framework

- A replacement for NumPy to use the power of GPUs
- A deep learning research platform that provides maximum flexibility and speed

Since it is highly embedded in Python, PyTorch is very Pythonic and easy-to-use

Pytorch Installation

Firstly check if your computer has discrete graphics card (GPU)



Install Nvidia driver: <https://zhuanlan.zhihu.com/p/54350088>

- CUDA 10.1
- cuDNN 7: [Installation guide](#)

Pytorch Installation

Select your configuration on this [website](#) and run the installation command

- Windows: Need to install Anaconda first
- WSL does not support GPU! Do NOT install Pytorch on WSL!
- Mac does not support GPU too (if you do not have external interface)!

e.g. For Windows with no GPUs

```
pip install torch==1.4.0+cpu torchvision==0.5.0+cpu -f https://download.pytorch.org/whl/torch_stable.html
```

Check if GPU works correctly by

```
import torch
print(torch.cuda.is_available())
```


Tutorials

PyTorch has very detailed documentations, make the best of them!

- Tutorials: <https://pytorch.org/tutorials/>
- Chinese tutorials: <https://pytorch.apachecn.org/>
- Documentation / API:
<https://pytorch.org/docs/stable/index.html>
- Deep Learning with PyTorch: A 60 Minute Blitz
 - Chinese version
 - You can download the .ipynb file or directly run on [Colab](#)

4

Summary

Summary

- Introduction
- Deep Learning Framework: PyTorch
 - Once you get into troubles concerning PyTorch, you can search [the Docs of PyTorch](#) for details. Alternatively you can try to find if there are similar problems on [PyTorch Discuss](#).
- Get through [cs231n](#)!

Assignment

Train you own network on CIFAR-10 and achieve 60%+ accuracy.

See `Assignments/PyTorch-CNN/main.ipynb` for more details.